Sravanth Kodavanti

+91-9550173710 ramasravanthkodavanti@gmail.com sravanthk27 Google Scholar OpenReview Website

About

Research Engineer specializing in **post-training optimization** and **efficient AI** for real-world impact. Skilled in advanced **inference optimization**, **model compression** (quantization), and **Neural Architecture Search**, delivering performance improvements in both industrial and research projects. Passionate about bridging cutting-edge AI methods with scalable solutions to address practical business needs.

Education

Indian Institute of Technology, Hyderabad **2020 - 2024**
Bachelor of Technology, **Major** - Computer Science and Engineering, **Minor** - Entrepreneurship CGPA - 8.58

Research Publications

1. Sowmya Vajrala*, Srinivas Miriyala*, ***Sravanth Kodavanti***
Towards Efficient Image Deblurring for Edge Deployment,
Under Review at ICASSP 2026

2. ***Sravanth Kodavanti****, Sowmya Vajrala*, Srinivas Miriyala*, Utsav Tiwari* *et al*
Unlocking the Edge Deployment and On-Device Acceleration of Multi-LoRA Enabled One-for-All Foundational LLM, *Under Review at EMNLP 2025 Industry Track*

3. ***Sravanth Kodavanti****, Srinivas Miriyala*, Sowmya Vajrala*, Vikram N R
On Distillation of Transformers into State-Space Models for Efficient Image Restoration,
Under Review at NeurIPS 2025

4. Srinivas Miriyala*, Sowmya Vajrala*, Hitesh Kumar, ***Sravanth Kodavanti***, Vikram N R
Mobile-friendly Image De-noising: Hardware-Conscious Optimization for Edge Application,
Proc. IEEE Int. Conf. Acoust., Speech, Signal Process. (ICASSP), 2025

Patents

1. ***Sravanth Kodavanti****, Srinivas Miriyala*, Sowmya Vajrala*
System and Method for Accelerating On-Device Large Language Model Inference.
Provisional Specification Filed. Proposes a novel self-speculative decoding technique to significantly improve efficiency and reduce latency in LLM inference on edge devices.

Work Experience

Samsung Research Institute Bangalore **Aug 2024 - Present**
Machine Learning Research Engineer ***Mentor: Dr. Srinivas Miriyala***

- Developing **Samsung Neural Acceleration Platform** for AI model acceleration and deployment on mobile devices, leveraging **Neural Architecture Search** (NAS) and **Quantization** for performance optimization.
- **Commercialization & Impact:**
  1. **Developed NAS and quantization-based models** for **low-light video de-noising** (**2.5**× faster; deployed on **Samsung A56**) and **demoire artifact removal** (**2.2**× faster; targeted for **Galaxy S26**), delivering advanced, commercialization-ready imaging solutions.
  2. **Accelerated on-device inference** for the **Samsung LLM Gauss - L, 3B Model** by implementing **Speculative Decoding**, achieving a **2X improvement** in tokens per second (toks/sec). Successfully integrated into the **Samsung S25 flagship series**.
- **Awards & Recognition:**
  1. **Spot Award (Q2 2025)** – For novel speculative decoding for LLM acceleration.
  2. **MD Project Incentive Award (2024–25)** – For AI model optimizations and deployment on edge devices.
  3. **Team Awesome Award (Q4 2024)** – For on-device optimization of **Samsung Gauss - L** in **S25** mobile series.

Stealth Startup - Gika.AI **June 2024 - Sep 2024**
AI Researcher ***Mentor: Dr. Manoj Aggarwal***

- Developed domain-specific search engines leveraging **Knowledge Graphs** (KG) to enhance semantic accuracy. Improved **Retrieval Augmented Generation** by using KG over vector databases for more precise, context-aware retrieval.
- Handled **Coreference** and **Entity Resolution** for fine-tuning data using LLMs such as **GPT-4, GPT-4o**, **LLAMA3**, **SpanBERT**, and **LingMess** with the **SpaCy** AI framework.

Hexagon R & D India **Jan 2024 - June 2024**
Machine Learning Intern ***Mentor: Ankan Sengupta***

- Developed a website for implementing segmentation & classification based on tags of various manufacturing plant sketches.Used various segmentation algorithms such as **RANSAC**, **DBSCAN**, **K-Means**.
- Used **Azure form recognizer** models & other OCR models for tag identification.
- Developed a **flask website** for text recognition on manufacturing plant sketches using **Form, Doc Recognizer** & also used OCR models for tag identification.
- Involved in an **LLM** research project. Compared the results for text generation between **Mixtral**, **Mistral**, **LLAMA2**.

OnePlus / Oppo (OPLUS) Mobiles India R & D **Jan 2023 - June 2023**
Research Intern - Device AI ***Mentor: Dr. C Shyam Anand***

- Implemented model compression techniques called **Quantization** , **Pruning** & **Distillation** on various deep learning models such as **ResNet** , **Yolo** , **ViTs** , **ConvNeXt** , **Stable Diffusion Models** & **Large Language Models** (LLAMA) . **Compressed all these models for deploying in mobile devices.**
- Involved in research work on **Neural Style Transfer** (NST).
- Impemented a DL model for **LaTeX - OCR** task . The model's aim is to detect & recognize the mathematical equations present in a research paper. Compressed the model for the integration with edge devices.

Teaching Experience

Department of Computer Science , IIT Hyderabad **Nov 2021 - Apr 2024**
Teaching Assistant

- Served as a **Teaching Assistant** for multiple courses, including:
  1. **Operating Systems** (CS3510) under ***Dr. Sathya Peri***
  2. **Discrete Mathematics** (CS1010) under ***Dr. Rakesh Vekat***
  3. **Database Management Systems** (CS3550) under ***Dr. Manish Singh***
  4. **Theory of Computation** (CS2030) under ***Dr. Subrahmanyam Kalyanasundaram***

Research Community Service

Research Reviewer **2025 – Present**

- Reviewed submissions for **ICML**, **NeurIPS**, **ACL**, **AAAI** & **ICASSP** workshops & main tracks.
- Provided constructive feedback to maintain academic rigor and contribute to the research community.

Projects

Continual Learning for 3D Point Cloud ***Guide: Dr. P K Srijith***

- Implemented Continual Learning on **Pointnet** architecture by addressing the challenge of catastrophic forgetting.
- Implemented **Knowledge Distillation** approach for the Continual Learning & Model was trained on **ModelNet10** dataset.

Computer Vision & NLP | Personal Projects

- Image denoising using **Auto - Encoders** . Model trained on **MNIST** dataset.
- Human Face generation using **GANs** . Model trained on **celeba** dataset.
- Underwater Object Detection and Classification using **DGYOLO**. Model trained on **URPC 2019** dataset.

Technical Skills

Languages: C, C++, Python, SQL, HTML, CSS
Frameworks & Tools : PyTorch, Tensorflow, ONNX, Hugging Face, Docker, Git, LaTeX, Markdown
Operating Systems : Linux, Windows
Familiar: Tensorrt, JavaScript

Achievements

- Secured AIR **1156** in Open Category & AIR **91** in EWS Category in IIT JEE Advanced 2020.
- Secured AIR **493** in Open Category & AIR **45** in EWS Category in JEE Main B Planning 2020.
- Secured **99.75** percentile in IIT JEE Main 2020.
- AP EAMCET 2020 Rank **135**.
- TS EAMCET 2020 Rank **307**.
- Solved around 700+ CP problems over multiple platforms LeetCode , CodeChef , CodeForces , GeeksForGeeks.

Leadership / Extracurricular

- Worked as a Core member for **Epoch : AI-ML club of IITH**.
- Worked as an **Internship** & **Placement** Coordinator for Office of Carrer Service of IITH.
- Member of IITH **Chess** team , Represented IITH at **InterIIT** & various competitions.